

NLP Course



Ctrl+K

EN

2,356

0. SETUP
1. TRANSFORMER MODELS
- Introduction
- Natural Language Processing
- Transformers, what can they do?
- How do Transformers work?
- Encoder models
- Decoder models
- Sequence-to-sequence models
- Bias and limitations
- Summary
- End-of-chapter quiz
2. USING TRANSFORMERS
3. FINE-TUNING A PRETRAINED MODEL
4. SHARING MODELS AND TOKENIZERS
5. THE DATASETS LIBRARY
6. THE TOKENIZERS LIBRARY
7. MAIN NLP TASKS
8. HOW TO ASK FOR HELP
9. BUILDING AND SHARING DEMOS NEW
- COURSE EVENTS

Join the Hugging Face community

and get access to the augmented documentation experience

Collaborate on models, datasets and Spaces

Faster examples with accelerated inference

Switch between documentation themes

Sign Up

 to get started

Summary

Ask a question

In this chapter, you saw how to approach different NLP tasks using the high-level `pipeline()` function from Transformers. You also saw how to search for and use models in the Hub, as well as how to use the Inference API to test the models directly in your browser.

We discussed how Transformer models work at a high level, and talked about the importance of transfer learning and fine-tuning. A key aspect is that you can use the full architecture or only the encoder or decoder, depending on what kind of task you aim to solve. The following table summarizes this:

Model	Examples	Tasks
Encoder	ALBERT, BERT, DistilBERT, ELECTRA, RoBERTa	Sentence classification, named entity recognition, extractive question answering
Decoder	CTRL, GPT, GPT-2, Transformer XL	Text generation
Encoder-decoder	BART, T5, Marian, mBART	Summarization, translation, generative question answering

